



Automated Machine Learning for Recommender System

Chen Gao

Ph.D Candidate, Tsinghua University

`gc16@mails.tsinghua.edu.cn`










Outline

- AutoML for Collaborative Filtering Task
- AutoML for Click-through Rate Prediction Task
- AutoML for Tuning Hyper-parameters in RecSys

Outline

- AutoML for Collaborative Filtering Task
- AutoML for Click-through Rate Prediction Task
- AutoML for Tuning Hyper-parameters in RecSys

Collaborative Filtering – Problem Setup

		M items				
						
N users		5		1		
				3		
			5			
				2		
						5

Study the fundamental CF problem [1]

- Data: a **rating matrix** with many unknown positions
- Task: estimate ratings on unknown positions
- Measurement: RMSE on estimated ratings.

More clarifications

- Take CF as a regression task here
- Implicit feedbacks [2,3] are **NOT** considered
- Side-information is [4] (e.g. user/item features) **NOT** assumed


[1]. Exact Matrix Completion via Convex Optimization. Foundations of Computational Mathematics. 2008

[2]. Collaborative filtering for implicit feedback datasets. ICDE 2008

[3]. Neural Collaborative Filtering. WWW 2017

[4]. Meta-Graph Based Recommendation Fusion over Heterogeneous Information Networks. KDD 2017 4

Collaborative Filtering – Low-rank approach [1,2]



- Partially observed rating matrix O can be well-approximated by a **low-rank** matrix X
- Matrix U : **user embedding**, matrix V : **item embedding**
- Rating prediction is given by an **inner product** of user embedding and item embedding

$$\min_{U,V} \sum_{(i,j) \in \Omega} \ell(\mathbf{u}_i^\top \mathbf{v}_j, \mathbf{O}_{ij}) + \underbrace{\frac{\lambda}{2} \|\mathbf{U}\|_F^2 + \frac{\lambda}{2} \|\mathbf{V}\|_F^2}_{\text{Regularization}}$$

Prediction
Observed Positions
Rating

[1]. Factorization meets the neighborhood: a multifaceted collaborative filtering model. KDD 2008

[2]. Exact Matrix Completion via Convex Optimization. Foundations of Computational Mathematics. 2008

Collaborative Filtering – Interaction Function (IFC)

$$\min_{\mathbf{U}, \mathbf{V}} \sum_{(i,j) \in \Omega} \ell(\mathbf{u}_i^\top \mathbf{v}_j, \mathbf{o}_{ij}) + \frac{\lambda}{2} \|\mathbf{U}\|_F^2 + \frac{\lambda}{2} \|\mathbf{V}\|_F^2,$$

1. Generate embedding vectors for users and items
2. Generate predictions by an **inner product** between embedding vectors
3. Evaluate predictions by a loss function on the training data set



Interaction function: how embedding vectors interact with each other

Low-rank approach: using inner product as interaction function

Collaborative Filtering – Alternative IFCs

Is low-rank approach good enough? **NO**, depends on tasks and datasets

- User **UA** likes item **IA** very much
- User **UB** likes item **IA** also very much



UA and UB should have very similar embeddings



Triangle inequality: $\|u_i - u_j\| \leq \|u_i - v_j\| + \|u_k - v_j\|$



Alternative IFC: **minus/plus** ^[1]

Failure of low-rank approach: $u_i^T v_j = u_k^T v_j$ does not mean u_i and u_k have small distance

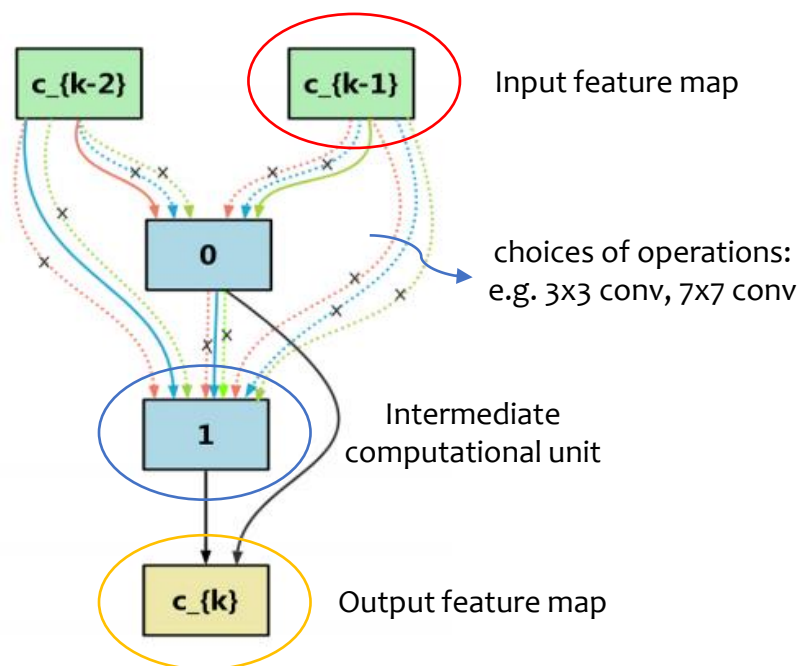
Collaborative Filtering – More Example IFCs

Table 1: Popular human-designed interaction functions (IFC) for CF, where H is a parameter to be trained. SIF searches a proper IFC from the validation set (i.e., by AutoML), while others are all designed by experts.

	IFC	operation	space	predict time	recent examples
human-designed	$\langle \mathbf{u}_i, \mathbf{v}_j \rangle$	inner product	$O((m+n)k)$	$O(k)$	MF [28], FM [37]
	$\mathbf{u}_i - \mathbf{v}_j$	plus (minus)	$O((m+n)k)$	$O(k)$	CML [19]
	$\max(\mathbf{u}_i, \mathbf{v}_j)$	max, min	$O((m+n)k)$	$O(k)$	ConvMF [25]
	$\sigma([\mathbf{u}_i; \mathbf{v}_j])$	concat	$O((m+n)k)$	$O(k)$	Deep&Wide [9]
	$\sigma(\mathbf{u}_i \odot \mathbf{v}_j + H[\mathbf{u}_i; \mathbf{v}_j])$	multi, concat	$O((m+n)k)$	$O(k^2)$	NCF [17]
	$\mathbf{u}_i * \mathbf{v}_j$	conv	$O((m+n)k)$	$O(k \log(k))$	ConvMF [25]
	$\mathbf{u}_i \otimes \mathbf{v}_j$	outer product	$O((m+n)k)$	$O(k^2)$	ConvNCF [16]
AutoML	SIF (proposed)	searched	$O((m+n)k)$	$O(k)$	—

Is there an absolute best IFC? : **NO**, depends on tasks and datasets ^[1]

NASP – Efficient NAS via Proximal Iterations^[1]



Supernet for CNN (a computational cell)

Edge representation:

(with **sparse constraint**)

$$\bar{O}(X) = \sum_{k=1}^d a_k O_k(X) \text{ where } \alpha \in C_1 \cap C_2,$$

$$C_1 = \{\alpha \mid \|\alpha\|_0 = 1\} \quad \text{and} \quad C_2 = \{\alpha \mid 0 \leq \alpha_k \leq 1\}.$$

Optimization Objective:

(**bi-level optimization**)

$$\min_{\alpha} \bar{\mathcal{L}}(w^*(\alpha), \alpha), \text{ s.t. } \begin{cases} w^*(\alpha) = \arg \min_w \mathcal{L}(w, \alpha) \\ \alpha \in C_1 \cap C_2 \end{cases} \quad \begin{matrix} \text{Validation set} \\ \text{Training set} \end{matrix}$$

α : architecture parameter, w : network parameter

Algorithm 1 Neural architecture search by proximal iterations (NASP) algorithm [47].

- 1: **require:** A mixture operation \bar{O} parametrized by (2), parameter w and step-size η ;
- 2: **while** not converged **do**
- 3: Obtain *discrete* architecture representation $\bar{\alpha} = \text{prox}_{C_1}(\alpha)$;
- 4: Update *continuous* architecture representation

$$\alpha = \text{prox}_{C_2}(\alpha - \nabla_{\bar{\alpha}} \bar{\mathcal{L}}(\bar{w}, \bar{\alpha}));$$

where $\bar{w} = w - \eta \nabla_w \mathcal{L}(w, \bar{\alpha})$ (is an approximation to $w^*(\bar{\alpha})$);

- 5: Get new *discrete* architecture $\bar{\alpha} = \text{prox}_{C_1}(\alpha)$;
- 6: Update w using $\nabla_w \mathcal{L}(w, \bar{\alpha})$ with $\bar{\alpha}$;
- 7: **end while**
- 8: **return** Searched architecture $\bar{\alpha}$.

A proximal algorithm for NAS

- **Space:** Super-net
- **Algorithm:** Proximal Gradient descent
- **Evaluation:** Parameter-sharing

10+ times faster than DARTS^[2]

[1] Efficient Neural Architecture Search via Proximal Iterations. AAAI 2020

[2] DARTS: Differentiable architecture search. ICLR 2019

Motivation

Is there an absolute best IFC? : **NO**, depends on tasks and datasets



Why not search (by **NAS**) an IFC from the data on the given task?

AutoML

- Search space
- Search Algorithm
- Evaluation Method

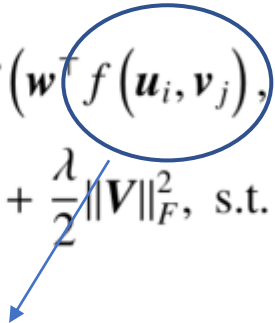


All these need to be carefully designed to be efficient and get better performance than existing IFCs.

HOW?

Search Space – Problem definition

Generalize standard CF objective:

$$\min F(\mathbf{U}, \mathbf{V}, \mathbf{w}) \equiv \sum_{(i,j) \in \Omega} \ell(\mathbf{w}^\top f(\mathbf{u}_i, \mathbf{v}_j), \mathbf{O}_{ij}) + \frac{\lambda}{2} \|\mathbf{U}\|_F^2 + \frac{\lambda}{2} \|\mathbf{V}\|_F^2, \text{ s.t. } \|\mathbf{w}\|_2 \leq 1,$$


Interaction Function:

- take user / item embeddings as input and output a vector

What can be a good candidate space?

DEFINITION 3.1 (AUTOML PROBLEM). Let \mathcal{M} be a performance measure (the lower the better) defined on the validation set $\bar{\Omega}$ (disjoint from Ω), and \mathcal{F} be a family of vector-valued functions with two vector inputs. The problem of searching for an interaction function (SIF), i.e., finding f^* , is defined as

$$f^* = \arg \min_{f \in \mathcal{F}} \sum_{(i,j) \in \bar{\Omega}} \mathcal{M}(f(\mathbf{u}_i^*, \mathbf{v}_j^*)^\top \mathbf{w}^*, \mathbf{O}_{ij}) \quad (6)$$
$$\text{s.t. } [\mathbf{U}^*, \mathbf{V}^*, \mathbf{w}^*] = \arg \min_{\mathbf{U}, \mathbf{V}, \mathbf{w}} F(\mathbf{U}, \mathbf{V}, \mathbf{w}),$$


where \mathbf{u}_i^* (resp. \mathbf{v}_j^*) is the i th column of \mathbf{U}^* (resp. j th column of \mathbf{V}^*).

- Too large: lead to extremely **huge computational** cost for subsequent optimization algorithms
- Too small: no need to AutoML, **worse performance** than existing IFCs

Search Space – Learning from Existing IFCs

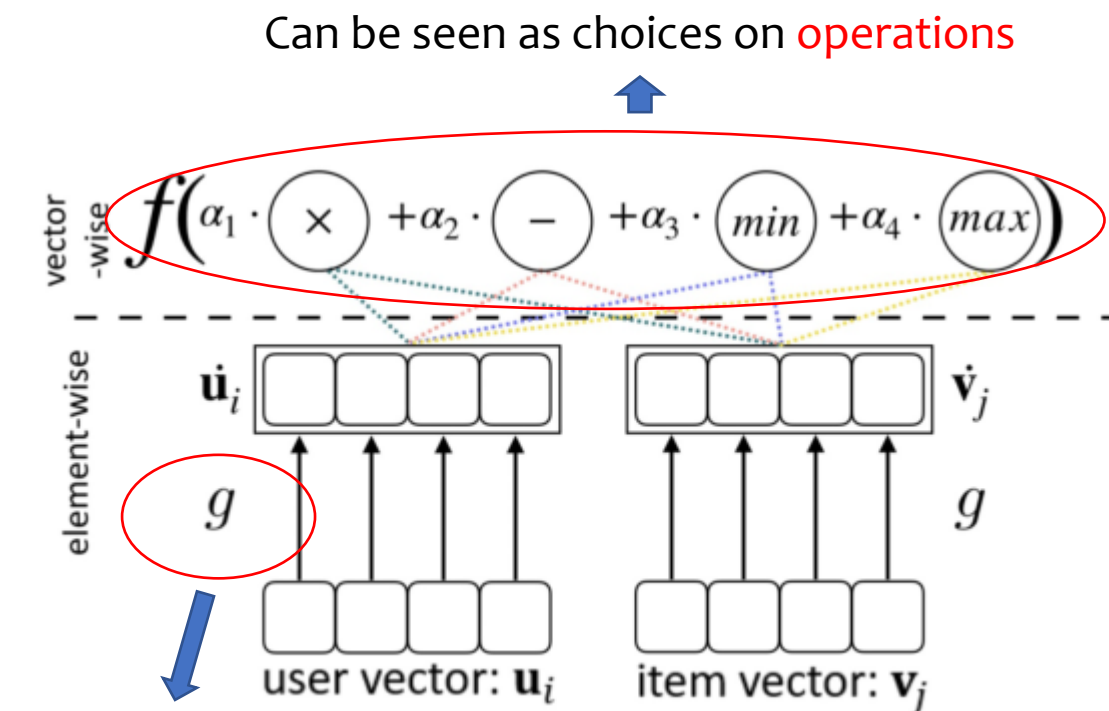
IFC	operation
$\langle \mathbf{u}_i, \mathbf{v}_j \rangle$	inner product
$\mathbf{u}_i - \mathbf{v}_j$	plus (minus)
$\max(\mathbf{u}_i, \mathbf{v}_j)$	max, min
$\sigma([\mathbf{u}_i; \mathbf{v}_j])$	concat
$\sigma(\mathbf{u}_i \odot \mathbf{v}_j + \mathbf{H}[\mathbf{u}_i; \mathbf{v}_j])$	multi, concat
$\mathbf{u}_i * \mathbf{v}_j$	conv
$\mathbf{u}_i \otimes \mathbf{v}_j$	outer product

Cut the search space into two blocks



- Vector-level: **simple** linear algebra operations
- Elementwise: **shared** nonlinear transformation

Bilevel Optimization Objective



Implement using a **small MLP**

A Supernet Representation

S : architecture hyper-parameters

T : parameters

$$\begin{aligned} \min_S \quad & H(S, T) \equiv \sum_{(i,j) \in \Omega} \mathcal{M}(h_\alpha(\mathbf{u}_i^*, \mathbf{v}_j^*)^\top \mathbf{w}_\alpha^*, \mathbf{O}_{ij}) \\ \text{s.t.} \quad & \alpha \in C \text{ and } T^* \equiv \{U^*, V^*, \{\mathbf{w}_m^*\}\} = \arg \min_T F_\alpha(T; S), \end{aligned} \quad (9) \quad \text{High level}$$

where F_α is the training objective:

$$\begin{aligned} F_\alpha(T; S) \equiv & \sum_{(i,j) \in \Omega} \ell(h_\alpha(\mathbf{u}_i, \mathbf{v}_j), \mathbf{O}_{ij}) + \frac{\lambda}{2} \|\mathbf{U}\|_F^2 + \frac{\lambda}{2} \|\mathbf{V}\|_F^2, \\ \text{s.t.} \quad & \|\mathbf{w}_m\|_2 \leq 1 \text{ for } m = 1, \dots, |O|. \end{aligned} \quad \text{Low level}$$

- High level: optimize S
- Low level: optimize T
- Bilevel programming is **expensive** to solve - T^* needs to be obtained from model training

Algorithm & Evaluation – Reusing NASP

Bilevel objective:

$$\begin{aligned} \min_S \quad & H(S, T) \equiv \sum_{(i,j) \in \Omega} \mathcal{M}(h_\alpha(\mathbf{u}_i^*, \mathbf{v}_j^*)^\top \mathbf{w}_\alpha^*, O_{ij}) \\ \text{s.t.} \quad & \alpha \in \mathcal{C} \text{ and } T^* \equiv \{U^*, V^*, \{\mathbf{w}_m^*\}\} = \arg \min_T F_\alpha(T; S), \end{aligned} \quad (9)$$

where F_α is the training objective:

$$\begin{aligned} F_\alpha(T; S) \equiv & \sum_{(i,j) \in \Omega} \ell(h_\alpha(\mathbf{u}_i, \mathbf{v}_j), O_{ij}) + \frac{\lambda}{2} \|\mathbf{U}\|_F^2 + \frac{\lambda}{2} \|\mathbf{V}\|_F^2, \\ \text{s.t.} \quad & \|\mathbf{w}_m\|_2 \leq 1 \text{ for } m = 1, \dots, |O|. \end{aligned}$$

Reuse NASP for fast optimization

- **Effectiveness:** Maintain discrete architectures for S
- **Efficiency:** Update both S and T in an **end-to-end and stochastic manner**

Algorithm 2 Searching Interaction Function (SIF) algorithm.

- 1: Search space \mathcal{F} represented by a structured MLP (Figure 1);
 - 2: **while** epoch $t = 1, \dots, T$ **do**
 - 3: Select one operation $\bar{\alpha} = \text{prox}_{C_1}(\alpha)$;
 - 4: *sample a mini-batch on validation data set*;
 - 5: Update continuous α for vector-wise operations


$$\alpha = \text{prox}_{C_2}(\alpha - \eta \nabla_{\bar{\alpha}} H(T, S));$$
 - 6: Update element-wise transformation

$$\mathbf{p} = \text{prox}_{\|\cdot\|_2 \leq 1}(\mathbf{p} - \eta \nabla_{\mathbf{p}} H(T, S)),$$

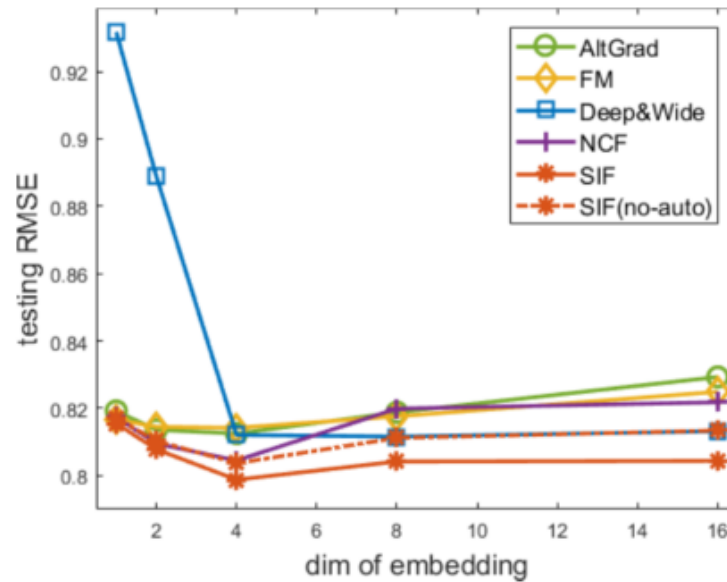
$$\mathbf{q} = \text{prox}_{\|\cdot\|_2 \leq 1}(\mathbf{q} - \eta \nabla_{\mathbf{q}} H(T, S));$$
 - 7: *sample a mini-batch on training data set*;
 - 8: Get selected operation $\bar{\alpha} = \text{prox}_{C_1}(\alpha)$;
 - 9: Update training parameters T with gradients on F_α ;
 - 10: **end while**
 - 11: **return** Searched interaction function (parameterized by α , \mathbf{p} and \mathbf{q} , see (7) and (8)).
-

Comparison with CF Approaches


(i) Alternating gradient descent (“**AltGrad**”); (ii) Factorization machine (“**FM**”); (iii) **Deep&Wide**; (iv) Neural collaborative filtering (“**NCF**”); (v) **SIF**; and (iv) **SIF(no-auto)**, architecture is optimized with training data



(a) MovieLens-100K.



(b) MovieLens-1M.




(c) Youtube.


Figure 2: Comparison of testing RMSEs between *SIF* and other CF approaches with different embedding dimension.

SIF is the best, and validation set helps architecture search


Comparison with CF Approaches



(a) MovieLens-100K.



(b) MovieLens-1M.




(c) Youtube.

Figure 3: Comparison of the convergence between *SIF* (with searched IFC) and other CF methods when embedded dimension is 8. *FM* and *HOFM* are not shown as their code donot support a callback to record testing performance.


Interaction function obtained from SIF can be trained as fast as state-of-the-art

Comparison with AutoML Approaches


(i) “Random”; (ii) “RL”: reinforcement learning; (iii) “Bayes”: HyperOpt



(a) MovieLens-100K.



(b) MovieLens-1M.

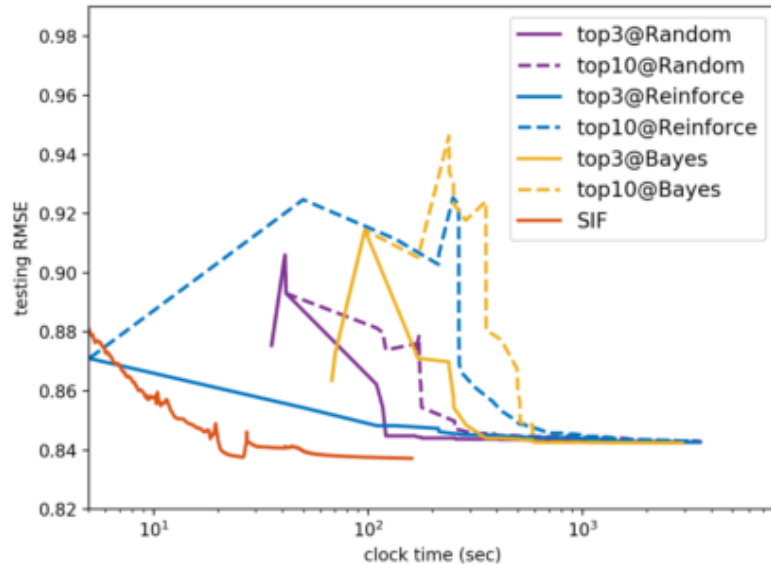


(c) Youtube.


Figure 4: Comparison of testing RMSEs between *SIF* and other AutoML approaches with different embedding dimensions. *Genapprox* is slow with bad performance, thus is not run on Youtube.

SIF can find better architecture than other AutoML search algorithms


Comparison with AutoML Approaches



(a) MovieLens-100K.



(b) MovieLens-1M.



(c) Youtube.

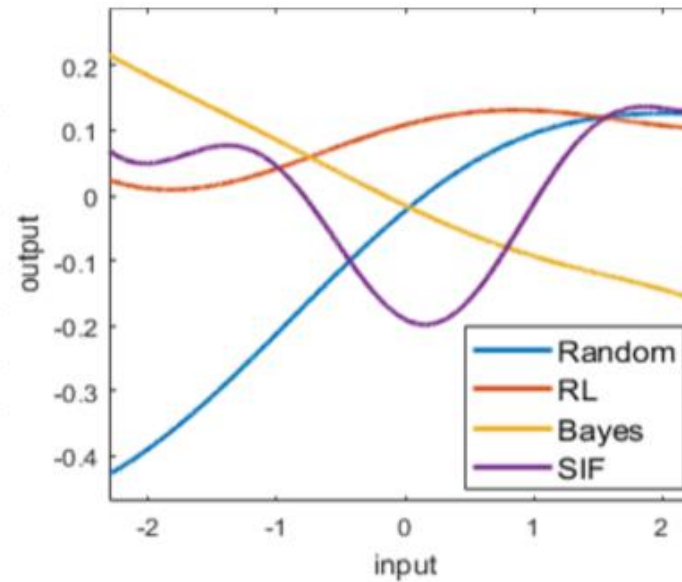
Figure 5: Comparison of search efficiency among *SIF* and other AutoML approaches when embedded dimension is 8.

SIF is much faster than other AutoML search algorithms

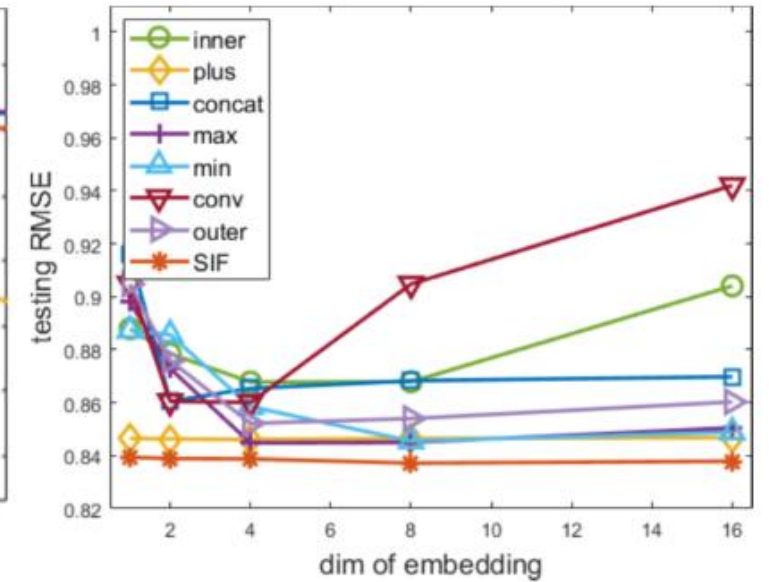
Case Study - Searched Interaction Functions (IFCs)

dim	Random	RL	Bayes	SIF
1	concat	concat	min	inner
2	concat	plus	concat	min
4	concat	concat	concat	min
8	plus	plus	plus	inner
16	concat	plus	plus	plus

(a) Operations (vector-wise).



(b) Non-linear transformation (element-wise).



(c) Performance of each single operation.

Figure 6: (a-b). Searched IFCs on MovienLens-100K with embedded dimension equals 8. (c). Performance comparison between SIF and each single operation on MovieLens-100K.

SIF can find more complex transformation and give better performance than any single operation

Ablation Study – Different search space





Figure 8: Comparison on different search space designs. Embedding dimension is 8.

NAS and Gen-approx are two general search space: inefficient to be searched

Outline

- AutoML for Collaborative Filtering Task
- AutoML for Click-through Rate Prediction Task
- AutoML for Tuning Hyper-parameters in RecSys

A pipeline of modern recommendation engine



Data Input : User-item interaction/rating data

Data Input : Rich attributes and context

Click-through rate prediction: **tabular data**

	age (n)	job (c)	marital (c)	education (c)	balance (n)	housing (c)
0	30	unemployed	married	primary	1787	no
1	33	services	married	secondary	4789	yes
2	35	management	single	tertiary	1350	yes
3	30	management	married	tertiary	1476	yes
4	59	blue-collar	married	secondary	0	yes
5	35	management	single	tertiary	747	no

An example of tabular data (UCI-Bank)

Cross-feature

- What is cross-features?
 - Taking cross-product of sparse features
- Why we need cross-features?
 - Capture the interaction among categorical features
 - Achieve great success in real-world business
- Traditional manner
 - LR+GBDT
 - DeepFM[1], xDeepFM[2], etc.

[1] Guo, Huifeng, et al. "DeepFM: a factorization-machine based neural network for CTR prediction." IJCAI 2017

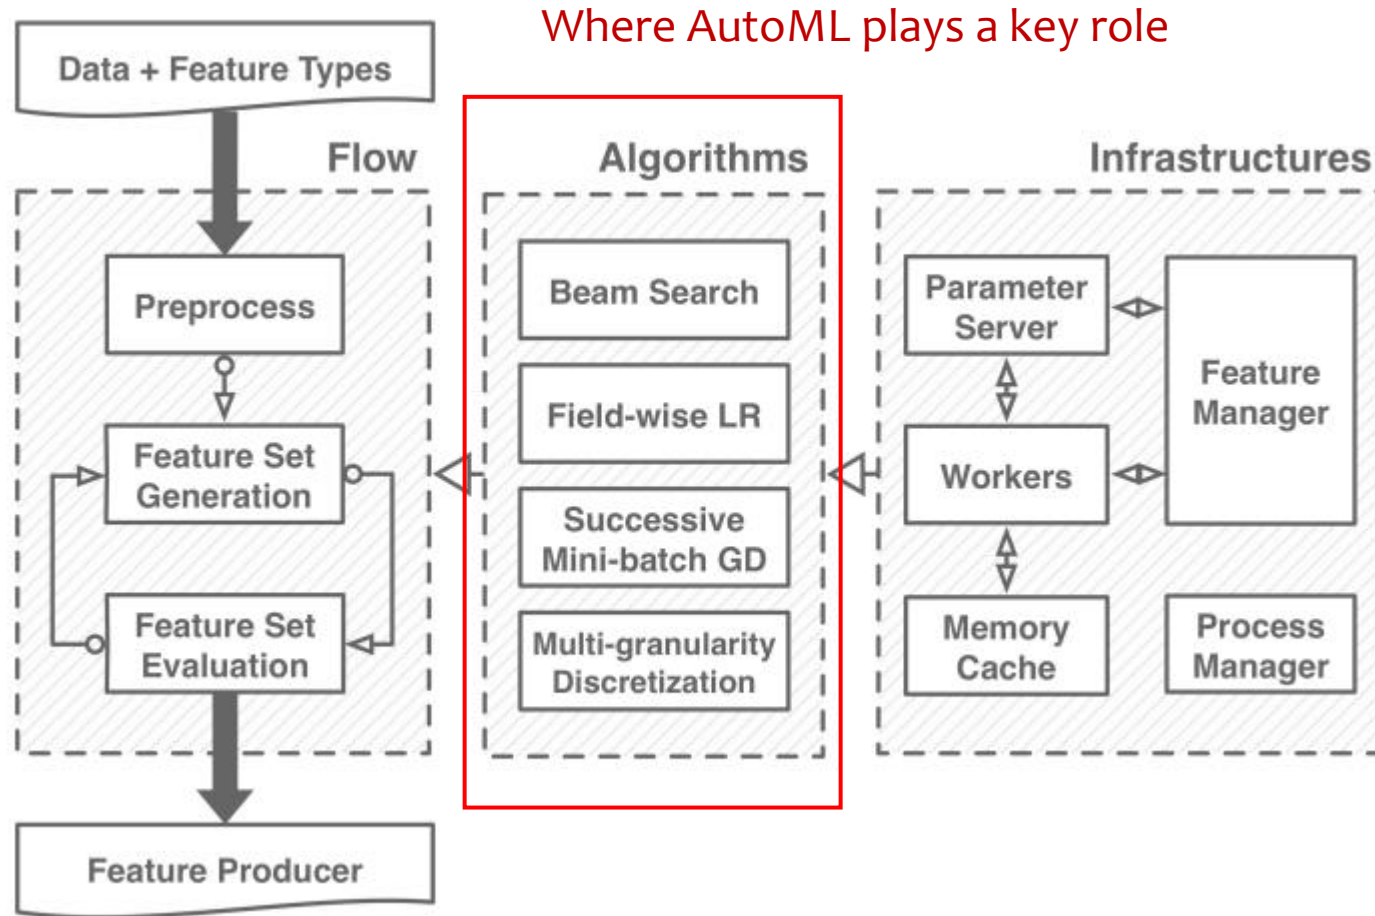
[2] Lian, Jianxun, et al. "xdeepfm: Combining explicit and implicit feature interactions for recommender systems." KDD 2018

Motivation

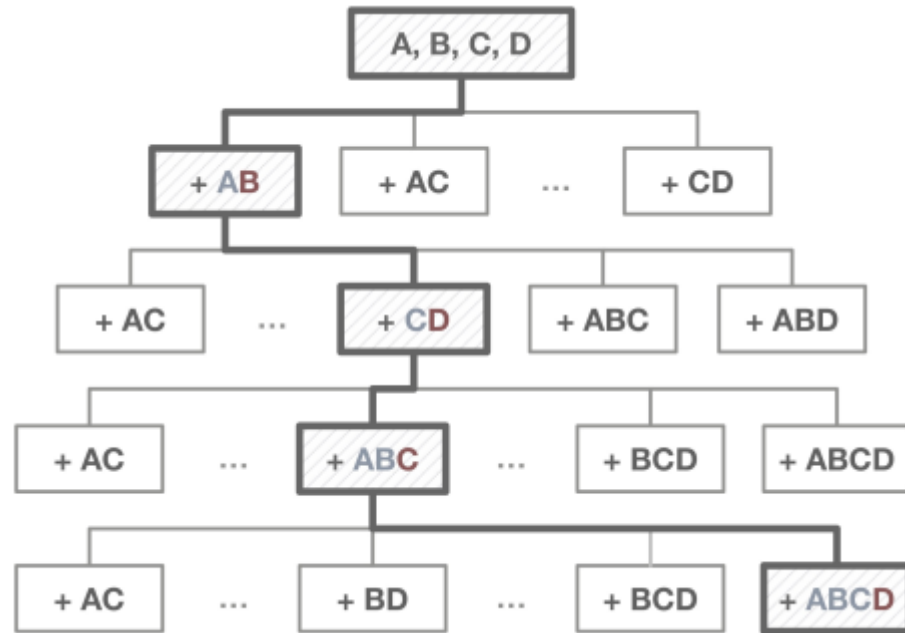
- Weaknesses of existing methods

Method	High-order Feature Crossing	Simplicity	Fast Inference	Interpretability
Search-based methods (e.g., [5, 34])	×	medium	√	√
Implicit deep-learning-based methods (e.g., [33, 42])	×	low	×	×
Explicit deep-learning-based methods (e.g., [26, 37])	×	low	×	√
AutoCross	√	high	√	√

Real-world System Framework of AutoCross



Method (feature search)



Search space and beam search strategy

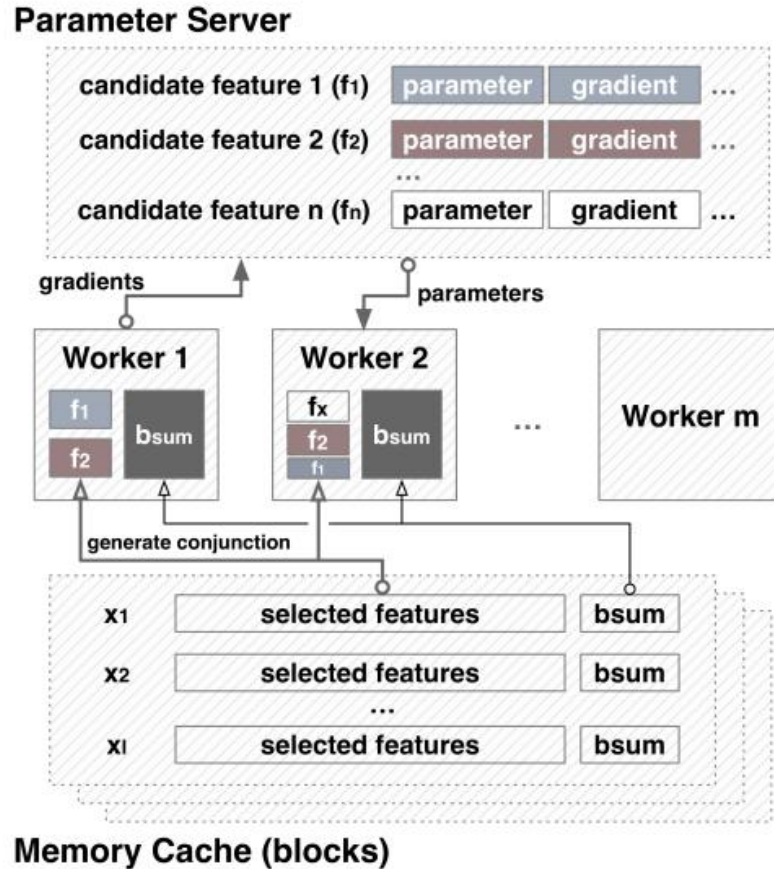
Algorithm 1 Feature Set Search Strategy in AutoCross.

Require: original feature set \mathcal{F} .

Ensure: solution S^* .

- 1: initialize current node $S^* \leftarrow \mathcal{F}$;
 - 2: **while** procedure not terminated **do**
 - 3: **Feature Set Generation:** expand S^* , generate its children node set $\text{children}(S^*)$ by adding to itself different pair-wise crossing of its elements;
 - 4: **Feature Set Evaluation:** evaluate all candidate feature sets in $\text{children}(S^*)$ and identify the best child S' ;
 - 5: visit S' : $S^* \leftarrow S'$
 - 6: **end while**
 - 7: **return** S^* .
-

Method (feature evaluation)



$$P(y = 1|\mathbf{x}) = s(\mathbf{w}^T \mathbf{x}) = s(\mathbf{w}_s^T \mathbf{x}_s + \mathbf{w}_c^T \mathbf{x}_c) = s(\mathbf{w}_c^T \mathbf{x}_c + b_{sum})$$

field-wise logistic regression (LR)

Algorithm 2 Successive Mini-batch Gradient Descent (SMBGD).

Require: set of candidate feature sets $\mathbb{S} = \{S_i\}_{i=1}^n$, training data equally divided into $N \geq \sum_{k=0}^{\lceil \log_2 n \rceil - 1} 2^k$ data blocks.

Ensure: best candidate S' .

- 1: **for** $k = 0, 1, \dots, \lceil \log_2 n \rceil - 1$ **do**
 - 2: use additional 2^k data blocks to update the field-wise LR models of all $S \in \mathbb{S}$, with warm-starting;
 - 3: evaluate the models of all S 's with validation AUC;
 - 4: keep the top half of candidates in \mathbb{S} : $\mathbb{S} \leftarrow \text{top_half}(\mathbb{S})$ (rounding down);
 - 5: break if \mathbb{S} contains only one element;
 - 6: **end for**
 - 7: **return** S' (the singleton element of \mathbb{S}).
-

field-wise logistic regression for feature evaluation based on PS architecture

Evaluations (Effectiveness)

Benchmark Datasets					
Method	Bank	Adult	Credit	Employee	Criteo
LR (base)	0.9400	0.9169	0.8292	0.8655	0.7855
AC+LR	0.9455	0.9280	0.8567	0.8942	0.8034
AC+W&D	0.9420	0.9260	0.8623	0.9033	0.8068
CMI+LR	0.9431	0.9153	0.8336	0.8901	0.7844
Deep	0.9418	0.9130	0.8369	0.8745	0.7985
xDeepFM	0.9419	0.9131	0.8358	0.8746	0.8059

Real-World Business Datasets					
Method	Data1	Data2	Data3	Data4	Data5
LR (base)	0.8368	0.8356	0.6960	0.6117	0.5992
AC+LR	0.8545	0.8536	0.7065	0.6276	0.6393
AC+W&D	0.8531	0.8552	0.7026	0.6260	0.6547
Deep	0.8479	0.8463	0.6936	0.6207	0.6509
xDeepFM	0.8504	0.8515	0.6936	0.6241	0.6514

Experimental results (test AUC) on benchmark and real-world business datasets.

AC+LR v.s. LR (base)					
Bank	Adult	Credit	Employee	Criteo	Average
0.585%	1.211%	3.316%	3.316%	2.279%	2.141%
Data1	Data2	Data3	Data4	Data5	Average
2.115%	2.154%	1.509%	2.599%	6.692%	3.014%

AC+W&D v.s. LR (base)					
Bank	Adult	Credit	Employee	Criteo	Average
0.213%	0.992%	3.992%	4.367%	2.712%	2.455%
Data1	Data2	Data3	Data4	Data5	Average
1.948%	2.346%	0.948%	2.338%	9.546%	3.368%

AC+W&D v.s. Deep					
Bank	Adult	Credit	Employee	Criteo	Average
0.021%	1.424%	3.035%	3.293%	1.039%	1.763%
Data1	Data2	Data3	Data4	Data5	Average
0.6133%	1.0516%	1.2976%	0.8539%	0.5361%	0.880%

Experimental results (test AUC) on benchmark and real-world business datasets.

Evaluations (Efficiency)

Benchmark Datasets				
Bank	Adult	Credit	Employee	Criteo
0.0267	0.0357	0.3144	0.0507	3.0817
Real-World Business Datasets				
Data1	Data2	Data3	Data4	Data5
0.9327	0.7973	1.5206	2.7572	5.1861

Cross feature generation time (unit: hour).

Benchmark Datasets					
Method	Bank	Adult	Credit	Employee	Criteo
AC+LR	0.00048	0.00048	0.00062	0.00073	0.00156
AC+W&D	0.01697	0.01493	0.00974	0.02807	0.02698
Deep	0.01413	0.01142	0.00726	0.02166	0.01941
xDeepFM	0.08828	0.05522	0.04466	0.06467	0.18985
Real-World Business Datasets					
Method	Data1	Data2	Data3	Data4	Data5
AC+LR	0.00367	0.00111	0.00185	0.00393	0.00279
AC+W&D	0.03537	0.01706	0.04042	0.02434	0.02582
Deep	0.02616	0.01348	0.03150	0.01414	0.01406
xDeepFM	0.32435	0.11415	0.40746	0.12467	0.13235

Inference latency comparison (unit: millisecond)


Outline

- AutoML for Collaborative Filtering Task
- AutoML for Click-through Rate Prediction Task
- AutoML for Tuning Hyper-parameters in RecSys



Regularization Tuning Headache



Bob



=




Model

$$l = \tilde{l}(\theta) + \lambda \underbrace{\|\theta\|^2}_{\text{Penalty}}$$

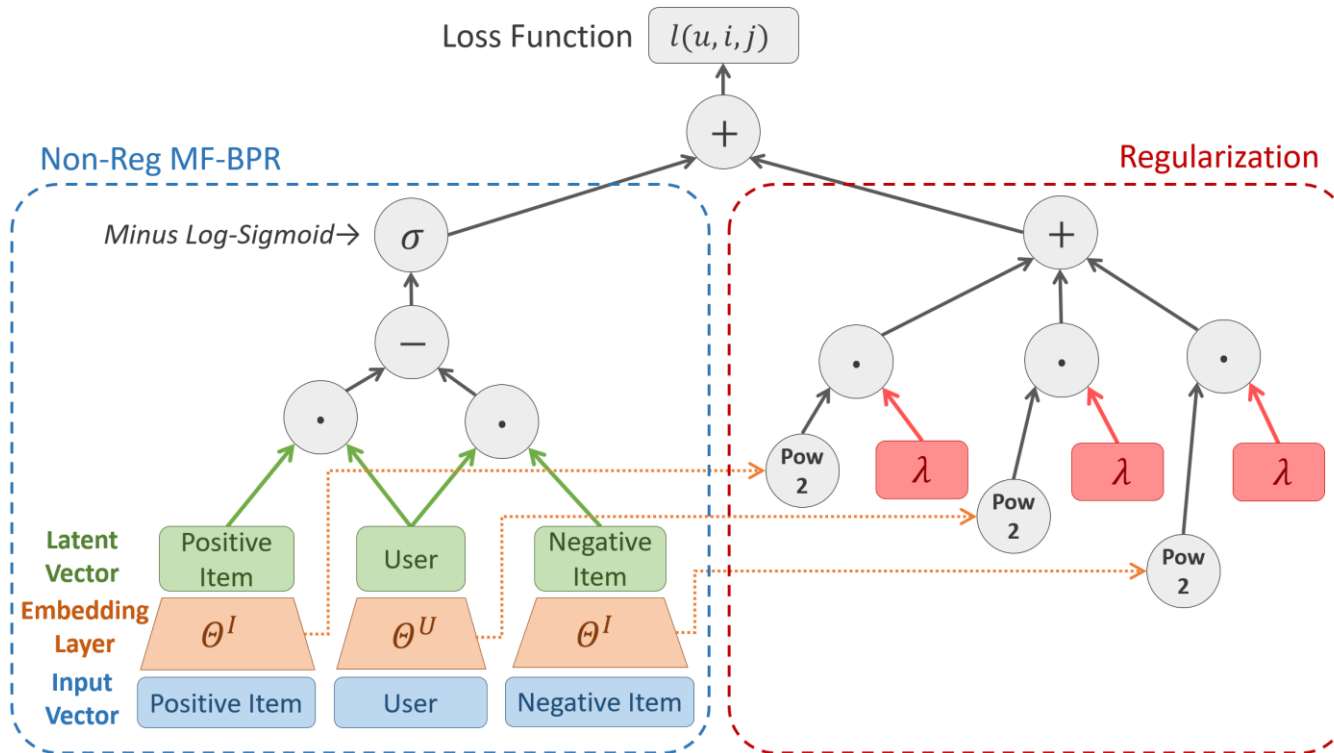
Penalty

Regularized Loss



What if we can
do the
regularization
automatically?

Matrix Factorization with Bayesian Personalized Ranking criterion




$$l_{S_T}(\Theta|\lambda) = \tilde{l}_{S_T}(\Theta) + \Omega(\Theta|\lambda)$$

$$= - \sum_{(u, i, j) \in S_T} \ln(\sigma(\hat{y}_{ui}(\Theta) - \hat{y}_{uj}(\Theta))) + \Omega(\Theta|\lambda)$$

S_T : training set,
 u : user,
 i : positive item,
 j : negative item,
 \hat{y}_{ui} : score function
 parametrized by MF for
 (u, i) pair
 \hat{y}_{uj} : score function
 parametrized by MF for
 (u, j) pair

Why hard to tune?

Hypothesis 1: fixed regularization strength throughout the process



Why hard to tune?

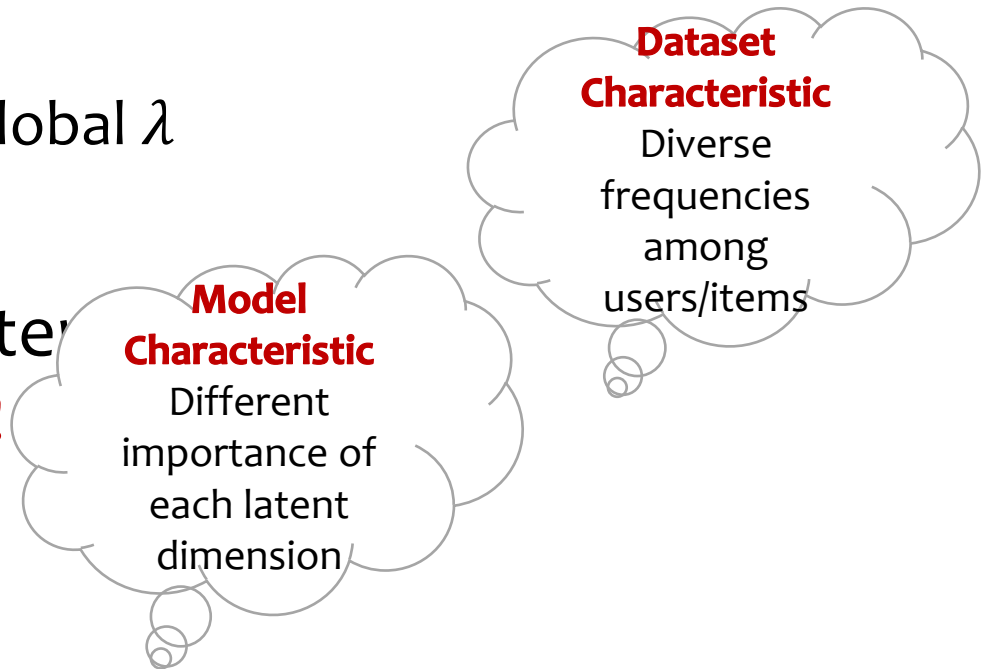
Hypothesis 2: compromise on regularization granularity

What we usually do to determine λ ?

- Usually Grid Search or Babysitting \rightarrow global λ

Fine-grained regularization works better

- But unaffordable if we use grid-search!
- Resort to automatic methods!



Alternating Optimization

$$\min_{\Lambda} \sum_{\{(u', i', j') \in S_V\}} l(u', i', j' | \arg \min_{\Theta} \sum_{\{(u, i, j) \in S_T\}} l(u, i, j | \Theta, \Lambda))$$

At iteration t

- Fix Λ , Optimize Θ

→ Conventional MF-BPR except λ is fine-grained now


Train the wheel!




- Fix Θ , Optimize Λ

→ Find Λ which achieve the smallest validation loss

Train the brake!



MF-BPR with fine-grained regularization



Fix Θ , Optimize Λ

Taking a greedy perspective, we look for Λ which can minimize the next-step validation loss

- If we keep using current Λ for next step, we would obtain $\bar{\Theta}_{t+1}$
- Given $\bar{\Theta}_{t+1}$, our aim is $\min_{\Lambda} l_{SV}(\bar{\Theta}_{t+1})$ with the constraint of non-negative Λ

But how to obtain $\bar{\Theta}_{t+1}$ without influencing the normal Θ update?

- Simulate* the MF update!
 - Obtain the gradients by combining the non-regularized part and penalty part

$$\frac{\partial \overline{l_{S_T}}}{\partial \Theta_t} = \frac{\partial \overline{\tilde{l}_{S_T}}}{\partial \Theta_t} + \frac{\partial \Omega}{\partial \Theta_t}$$


Λ is the only variable here


- Simulate the operations that the MF optimizer would take

$$\bar{\Theta}_{t+1} = f\left(\Theta_t, \frac{\partial \overline{l_{S_T}}}{\partial \Theta_t}\right)$$


f denotes the MF update function


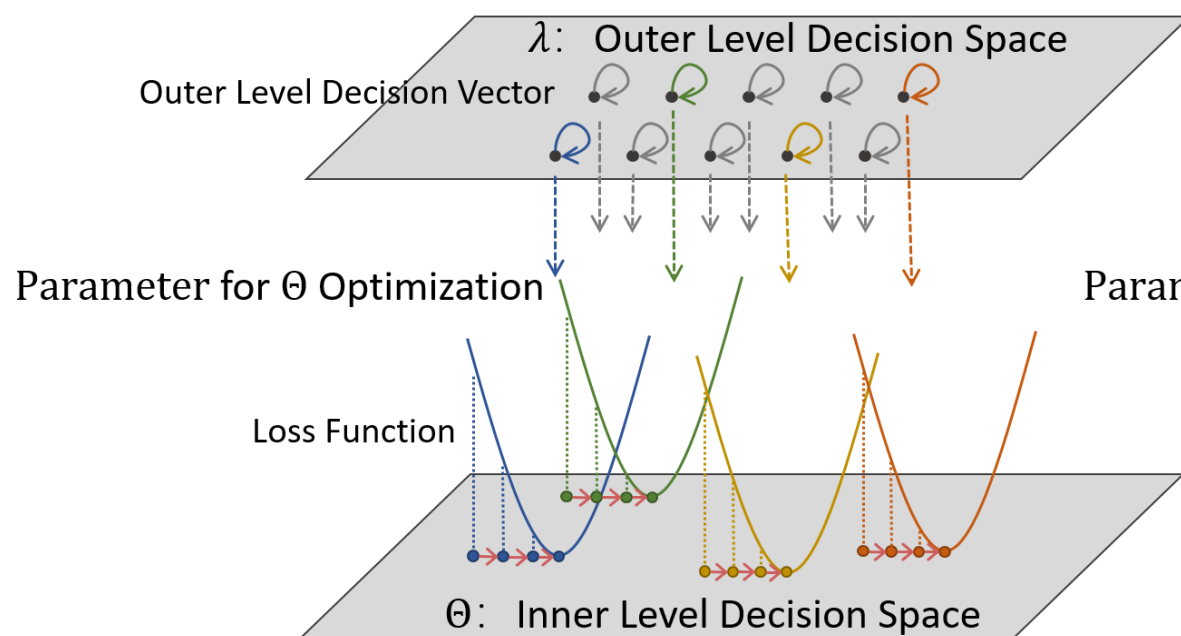
*: Using $\bar{}$ over the letters to distinguish the simulated ones with normal ones

Fix Θ , Optimize Λ in Auto-Differentiation



Another Perspective on Regularization Tuning

Λ – trajectory



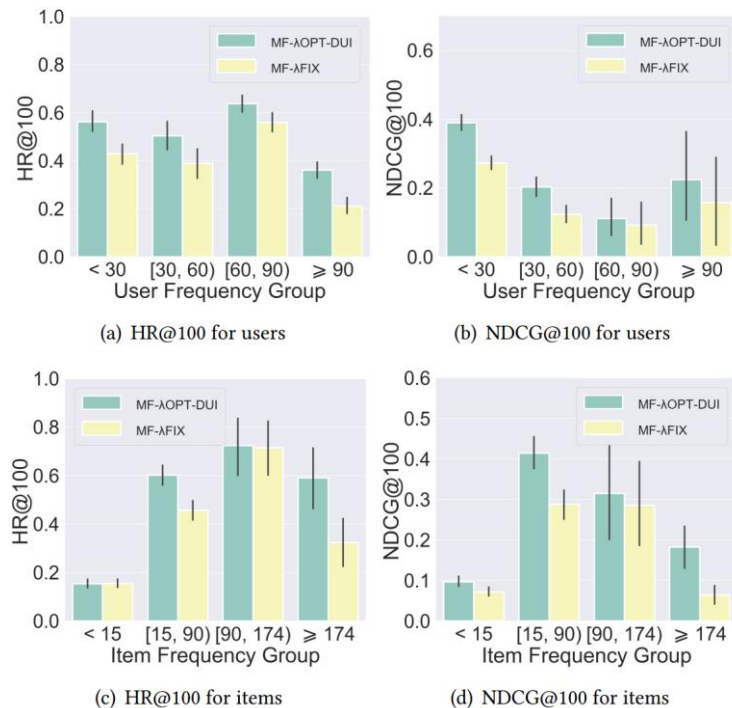
Result #1 Performance Comparison

Method	Amazon Food Review					MovieLens 10M				
	AUC	HR@50	HR@100	NDCG@50	NDCG@100	AUC	HR@50	HR@100	NDCG@50	NDCG@100
SGDA [26]	0.8130	0.1786	0.3857	0.1002	0.1413	0.9497	0.2401	0.3706	0.0715	0.0934
AMF [15]	0.8197	0.3541	0.4200	0.2646	0.2552	0.9495	0.2625	0.3847	0.0787	0.0985
NeuMF [16]	0.8103	0.3537	0.4127	0.2481	0.2218	0.9435	0.2524	0.3507	0.0760	0.0865
MF- λ Fix	0.8052	0.3482	0.4163	0.2251	0.2217	0.9497	0.2487	0.3779	0.0727	0.0943
MF- λ Opt -D	0.8109	0.2134	0.3910	0.1292	0.1543	0.9501	0.2365	0.3556	0.0715	0.0909
-DU	0.8200	0.3694	0.4814	0.2049	0.2570	0.9554	0.2743	0.4109	0.0809	0.1031
-DI	0.8501	0.2966	0.4476	0.1642	0.2039	0.9516	0.2648	0.3952	0.0804	0.1013
-DUI	0.8743	0.4470	0.5251	0.2946	0.2920	0.9575	0.3027	0.4367	0.0942	0.1158

1. **Overall:** MF- λ Opt-DUI achieves the best performance, demonstrating the effect of fine-grained adaptive regularization. (approx. 10%-20% gain over baselines)
2. **Dataset:** Performance improvement on *Amazon Food Review* is larger than that on *MovieLens 10M*. This might due to the dataset size and density. *Amazon Food Review* has a smaller number of interactions. Complex models like NeuMF or AMF wouldn't be at their best condition. Also, smart regularization is necessary for different users/items, explaining why SGDA and MF- λ Opt-DUI performs worse. In our experiments, we also observe more fluctuation of training curves on *Amazon Food Review* for the adaptive λ methods.
3. **Variants of regularization granularity:** Although MF- λ Opt-DUI consistently performs best, MF- λ Opt-DU/ or MF- λ Opt-DU doesn't provide as much gain over the baselines, which might be due to merely addressing the regularization for **partial** model parameters.

Result #2: Sparseness & Activeness

Does the performance improvement come from addressing different users/items?




Group users/items according to their frequencies and check the recommendation performance of each group, using *Amazon Food Review* as an example; black line indicates variance


1. **User with varied frequencies:** For users, MF-λOpt-DUI lifts HR@100 and NDCG@100. Compared to global MF-λOpt-DUI, fine-grained regularization addressing users of different frequencies better.
2. **Item with varied frequencies:** For items, similar lift can be observed except that only slight lift for HR@100 of the <15 group and [90, 174) group.
3. **Variance within the same group:** Although the average lift can be observed across groups, the variance demonstrate that there are factors other than frequency which influence the recommendation performance.

Result #3: Analysis of λ -trajectory

How does MF- λ Opt-DUI address different users/items?



(a) For users on Amazon Food Review



For each user/item, we cache the λ from Epoch 0 to Epoch 3200 (almost converged). λ s of users/items with the same frequency are averaged. The darker colors indicates larger λ .

1. **λ vs. user frequency:** At the same training stage, Users with higher frequencies are allocated larger λ . Active users have more data and the model learns from the data so quickly that it might get overfitting to them, making strong regularization necessary. A global λ , either small or large, would fail to satisfy both active users and sparse users.
2. **λ vs. item frequency:** Similar as the analysis of users though not so obvious. Items with higher frequencies are allocated larger λ .
3. **λ vs. training progress:** As training goes on, λ s gets larger gradually. Hence stronger regularization strengths are enforced at the late stage of training while the model is allowed to learn sufficiently at the beginning.

Conclusion

- AutoML can
 - help choosing models
 - select or generate data/feature
 - and even help tune hyper-parameters.